

Part_I_exploration_template

October 24, 2022

1 Part I - Fordgo Bike Trip Dataset

1.1 by Matthew Chan

2 Introduction

The following dataset is about the ford gobike dataset, it contains the trips of its app users for the month of February.

3 Preliminary Wrangling

```
In [1]: # import all packages and set plots to be embedded inline
```

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
%matplotlib inline
```

```
In [2]: # loading the dataset
```

```
fordbike = pd.read_csv('201902-fordgobike-tripdata.csv')
```

```
In [3]: # defining a function to explore and summarize datasets
```

```
def summary(df: pd.DataFrame) -> pd.DataFrame:
    """
    Returns the data type, # of null rows and unique rows of a given dataframe

    Args:
        A dataframe of n_rows and n_columns
    Returns:
        A dataframe of 4 columns [data_type, non-null-rows, unique_rows, missing_rows]
    """
    concat = pd.concat(
        [df.dtypes.to_frame(), df.count().to_frame(), df.nunique(), df.isnull().sum()],
        )
    concat.columns = ['data_type', 'non-null-rows', 'unique_rows', 'missing_rows']
    return concat
```

```
In [4]: # creating a new copy so we can revert to original data if needed after all the cleaning
bike_df = fordbike.copy()

# summarizing the dataset
summary(bike_df)
```

```
Out[4]:
```

	data_type	non-null-rows	unique_rows	missing_rows
duration_sec	int64	183412	4752	0
start_time	object	183412	183401	0
end_time	object	183412	183397	0
start_station_id	float64	183215	329	197
start_station_name	object	183215	329	197
start_station_latitude	float64	183412	334	0
start_station_longitude	float64	183412	335	0
end_station_id	float64	183215	329	197
end_station_name	object	183215	329	197
end_station_latitude	float64	183412	335	0
end_station_longitude	float64	183412	335	0
bike_id	int64	183412	4646	0
user_type	object	183412	2	0
member_birth_year	float64	175147	75	8265
member_gender	object	175147	3	8265
bike_share_for_all_trip	object	183412	2	0

3.1 Cleaning and transforming the Dataset

We are going to perform the following data cleaning/preparations on the dataset: 1. Assigning the correct/appropriate datatypes 2. Adding an 'age' column to denote users age 3. Editing the datetime format of certain columns for readability and easier analysis

3.1.1 Casting into appropriate datatypes

```
In [5]: # casting the following variables into their correct/appropriate datatypes
bike_df = (
    bike_df
    .astype({'start_time': 'datetime64[ns]',
             'end_time': 'datetime64[ns]',
             'start_station_id': 'str',
             'end_station_id': 'str',
             'bike_id': 'str',
             'member_gender': 'category',
             'user_type': 'category'})
)
```

3.1.2 Creating new columns based on the start_time

```
In [6]: # adding age of riders 'age' column into dataset
# adding new columns - 'date, hour, day and month' for datetime formats
```

```

bike_df = (
    bike_df
    .assign(age=lambda a: 2022-a.member_birth_year,
            start_date=lambda sd: sd.start_time.dt.strftime('%Y-%m-%d'),
            start_hour=lambda x: x.start_time.dt.strftime('%H'),
            start_day=lambda x: x.start_time.dt.strftime('%A'),
            start_month=lambda x: x.start_time.dt.strftime('%B'))
    )

summary(bike_df)

```

```

Out[6]:

```

	data_type	non-null-rows	unique_rows	\
duration_sec	int64	183412	4752	
start_time	datetime64[ns]	183412	183401	
end_time	datetime64[ns]	183412	183397	
start_station_id	object	183412	330	
start_station_name	object	183215	329	
start_station_latitude	float64	183412	334	
start_station_longitude	float64	183412	335	
end_station_id	object	183412	330	
end_station_name	object	183215	329	
end_station_latitude	float64	183412	335	
end_station_longitude	float64	183412	335	
bike_id	object	183412	4646	
user_type	category	183412	2	
member_birth_year	float64	175147	75	
member_gender	category	175147	3	
bike_share_for_all_trip	object	183412	2	
age	float64	175147	75	
start_date	object	183412	28	
start_hour	object	183412	24	
start_day	object	183412	7	
start_month	object	183412	1	

	missing_rows
duration_sec	0
start_time	0
end_time	0
start_station_id	0
start_station_name	197
start_station_latitude	0
start_station_longitude	0
end_station_id	0
end_station_name	197
end_station_latitude	0
end_station_longitude	0
bike_id	0
user_type	0

member_birth_year	8265
member_gender	8265
bike_share_for_all_trip	0
age	8265
start_date	0
start_hour	0
start_day	0
start_month	0

In [7]: # seems like there are missing rows in our age dataset, we will remove these columns
bike_df[bike_df['age'].isnull()].sample(5)

```
Out[7]:
```

	duration_sec		start_time		end_time	\
67099	684	2019-02-20	08:11:59.918	2019-02-20	08:23:24.379	
141146	6010	2019-02-07	16:44:20.218	2019-02-07	18:24:30.523	
35854	123	2019-02-23	17:40:47.149	2019-02-23	17:42:50.315	
173009	179	2019-02-03	13:55:55.643	2019-02-03	13:58:54.752	
127325	167	2019-02-10	16:15:53.516	2019-02-10	16:18:40.527	

	start_station_id		start_station_name	\
67099	67.0	San Francisco Caltrain Station 2	(Townsend St...	
141146	13.0		Commercial St at Montgomery St	
35854	233.0	4th Ave at E 12th St	(Temporary Location)	
173009	173.0		Shattuck Ave at 55th St	
127325	345.0		Hubbell St at 16th St	

	start_station_latitude	start_station_longitude	end_station_id	\
67099	37.776639	-122.395526	15.0	
141146	37.794231	-122.402923	368.0	
35854	37.795913	-122.255547	200.0	
173009	37.840364	-122.264488	169.0	
127325	37.766483	-122.398279	114.0	

	end_station_name	\
67099	San Francisco Ferry Building (Harry Bridges Pl...	
141146	Myrtle St at Polk St	
35854	2nd Ave at E 18th St	
173009	Bushrod Park	
127325	Rhode Island St at 17th St	

	end_station_latitude	...	bike_id	user_type	\
67099	37.795392	...	48	Subscriber	
141146	37.785434	...	2351	Customer	
35854	37.800214	...	4682	Subscriber	
173009	37.846516	...	662	Customer	
127325	37.764478	...	5522	Subscriber	

	member_birth_year	member_gender	bike_share_for_all_trip	age	\
--	-------------------	---------------	-------------------------	-----	---

67099	NaN	NaN	No NaN
141146	NaN	NaN	No NaN
35854	NaN	NaN	No NaN
173009	NaN	NaN	No NaN
127325	NaN	NaN	No NaN

	start_date	start_hour	start_day	start_month
67099	2019-02-20	08	Wednesday	February
141146	2019-02-07	16	Thursday	February
35854	2019-02-23	17	Saturday	February
173009	2019-02-03	13	Sunday	February
127325	2019-02-10	16	Sunday	February

[5 rows x 21 columns]

```
In [8]: # removing null/na numbers of members without age
bike_df = bike_df.dropna(subset=['age'])
# casting age into appropriate datatype (int)
bike_df = bike_df.astype({'age': 'int32'})

summary(bike_df)
```

```
Out[8]:
```

	data_type	non-null-rows	unique_rows \
duration_sec	int64	175147	4432
start_time	datetime64[ns]	175147	175136
end_time	datetime64[ns]	175147	175134
start_station_id	object	175147	330
start_station_name	object	174952	329
start_station_latitude	float64	175147	334
start_station_longitude	float64	175147	335
end_station_id	object	175147	330
end_station_name	object	174952	329
end_station_latitude	float64	175147	335
end_station_longitude	float64	175147	335
bike_id	object	175147	4635
user_type	category	175147	2
member_birth_year	float64	175147	75
member_gender	category	175147	3
bike_share_for_all_trip	object	175147	2
age	int32	175147	75
start_date	object	175147	28
start_hour	object	175147	24
start_day	object	175147	7
start_month	object	175147	1

	missing_rows
duration_sec	0
start_time	0

end_time	0
start_station_id	0
start_station_name	195
start_station_latitude	0
start_station_longitude	0
end_station_id	0
end_station_name	195
end_station_latitude	0
end_station_longitude	0
bike_id	0
user_type	0
member_birth_year	0
member_gender	0
bike_share_for_all_trip	0
age	0
start_date	0
start_hour	0
start_day	0
start_month	0

3.1.3 Structure of the dataset

Each dataset represents a unique, single trip by a user in the month of February. It contains the start, end, duration and other spatial data of the trip. There are a total of 16 variables, with around 183k records. These records can be broadly categorized into 3 main groups:

1. Timeseries data - e.g, duration_sec, start/end_time, birth_year
2. Spatial data - e.g, start/end_station_longitude/latitude
3. Member information - gender, age
4. Derived features - start_date/hour/day

3.1.4 What is/are the main feature(s) of interest in your dataset?

We are interesting in understanding the usage patterns of riders, which involves their ride characteristics (e.g, trip durations) alongside the profile of a user (e.g, age, gender, user_type)

3.1.5 What features in the dataset do you think will help support your investigation into your feature(s) of interest?

The start and end time variables will be useful in understanding the usage patterns. Also, start_day/month allows us to understand which days are most popular with users. The member information like user type, gender and age will help us find out who are the main target customer groups, use the different groups to summarize bike usage data to see if there is any special pattern associated with a specific group of riders.

```
In [9]: # overview of dataset
        bike_df.sample(5)
```

```
Out[9]:
```

	duration_sec		start_time		end_time	\
18613	618	2019-02-26	18:10:51.600	2019-02-26	18:21:10.171	
135259	472	2019-02-08	14:16:20.930	2019-02-08	14:24:13.784	
47066	535	2019-02-22	09:04:38.510	2019-02-22	09:13:33.526	
42025	1716	2019-02-22	17:35:46.618	2019-02-22	18:04:23.168	
124256	479	2019-02-11	08:42:07.088	2019-02-11	08:50:06.996	

	start_station_id		start_station_name	\
18613	81.0		Berry St at 4th St	
135259	27.0		Beale St at Harrison St	
47066	130.0		22nd St Caltrain Station	
42025	8.0		The Embarcadero at Vallejo St	
124256	31.0		Raymond Kimbell Playground	

	start_station_latitude	start_station_longitude	end_station_id	\
18613	37.775880	-122.393170	126.0	
135259	37.788059	-122.391865	81.0	
47066	37.757288	-122.392051	30.0	
42025	37.799953	-122.398525	77.0	
124256	37.783813	-122.434559	19.0	

		end_station_name	end_station_latitude	\
18613		Esprit Park	37.761634	
135259		Berry St at 4th St	37.775880	
47066	San Francisco Caltrain (Townsend St at 4th St)		37.776598	
42025		11th St at Natoma St	37.773507	
124256		Post St at Kearny St	37.788975	

	...	bike_id	user_type	member_birth_year	member_gender	\
18613	...	567	Customer	1991.0	Male	
135259	...	5529	Subscriber	1982.0	Male	
47066	...	4332	Subscriber	1983.0	Female	
42025	...	5242	Subscriber	1952.0	Female	
124256	...	4848	Subscriber	1994.0	Male	

	bike_share_for_all_trip	age	start_date	start_hour	start_day	\
18613	No	31	2019-02-26	18	Tuesday	
135259	No	40	2019-02-08	14	Friday	
47066	No	39	2019-02-22	09	Friday	
42025	No	70	2019-02-22	17	Friday	
124256	No	28	2019-02-11	08	Monday	

	start_month
18613	February
135259	February

```
47066    February
42025    February
124256    February
```

```
[5 rows x 21 columns]
```

4 Univariate Exploration

```
In [97]: # define a default chart size for all visualizations
plt.rcParams["figure.figsize"] = (9,7)
# seaborn plot size
plt.figure(figsize=[9, 7])
```


```
Out[97]: <matplotlib.figure.Figure at 0x7f90fc839e10>
```

```
<matplotlib.figure.Figure at 0x7f90fc839e10>
```

4.1 Visualization 1: What is the distribution of rides by user_type?

```
In [11]: # segmenting out our dataset for the pie plot
bike_df.user_type.value_counts().plot(kind='pie', autopct='%1.1f%%', shadow=True, color
# naming the line chart
plt.title("Distribution of rides by usertype", fontsize=17)
```

```
Out[11]: Text(0.5,1,'Distribution of rides by usertype')
```



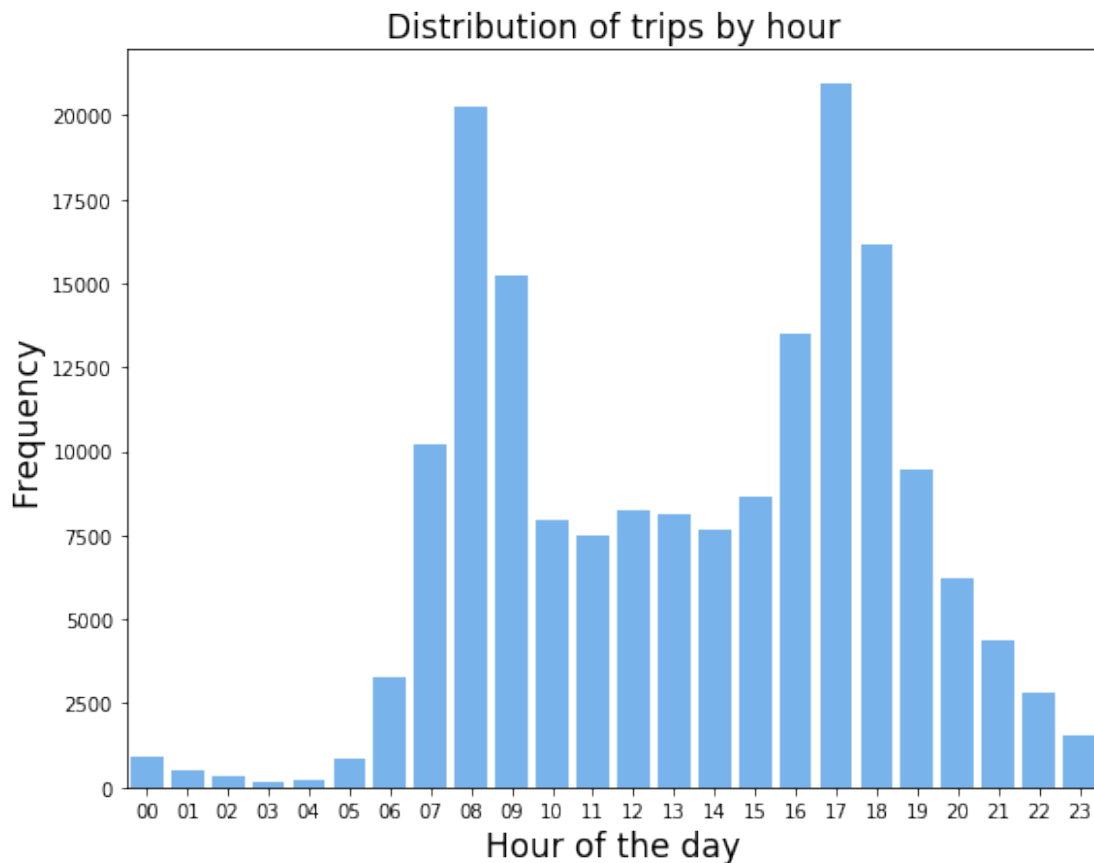
As observed, a significant majority of rides were completed by subscribers (90.5%) as compared to customers at 9.5%.

4.2 Visualization 2: What is the trip distribution over the hours of a day?

```
In [28]: trip_hour_distribution = sns.countplot(data=bike_df, x='start_hour', color='#66b3ff')

# set title of plot
trip_hour_distribution.axes.set_title('Distribution of trips by hour', fontsize=17)
# set x_label of plot
trip_hour_distribution.axes.set_xlabel('Hour of the day', fontsize=17)
# set y_label of plot
trip_hour_distribution.axes.set_ylabel('Frequency', fontsize=17)
```

```
Out[28]: Text(0,0.5, 'Frequency')
```



Trips peaked around 8am and 5pm, which seems to be during the rush hour periods.


4.3 Visualization 3: What is the trip distribution over days of the week?

```
In [171]: # defining days of the week for our chart
days_of_week = ['Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday', 'Saturday', 'Sunday']
# declaring the day list as ordered data
days_of_week_cat = pd.api.types.CategoricalDtype(ordered=True, categories=days_of_week)
bike_df = bike_df.astype({'start_day': days_of_week_cat})

# defining the countplot
trip_day_distribution = sns.countplot(data=bike_df, x='start_day', color='#66b3ff')

# set title of plot
trip_day_distribution.axes.set_title('Distribution of trips by day of week', fontsize=14)
# set x_label of plot
trip_day_distribution.axes.set_xlabel('Day of the week', fontsize=17)
# set y_label of plot
trip_day_distribution.axes.set_ylabel('Frequency', fontsize=17)

Out[171]: Text(0,0.5,'Frequency')
```




Trips typically happened more during the weekdays rather than the weekends.

4.4 Visualization 4: What is the trip distribution of days in the month of February?

```
In [91]: # defining days of the month for our chart
days_of_month = bike_df.start_date.unique().tolist().reverse()
days_of_month_cat = pd.api.types.CategoricalDtype(ordered=True, categories=days_of_month)
bike_df = bike_df.astype({'start_date': days_of_month_cat})

# defining the countplot
trip_month_distribution = sns.countplot(data=bike_df, x='start_date', color='#66b3ff')
# rotation x-axis for clarity
plt.xticks(rotation=45)
# set title of plot
trip_month_distribution.axes.set_title('Distribution of trips by day in February', font
# set x_label of plot
trip_month_distribution.axes.set_xlabel('Day of the Month', fontsize=17)
# set y_label of plot
trip_month_distribution.axes.set_ylabel('Frequency', fontsize=17)
```

```
Out[91]: Text(0,0.5,'Frequency')
```

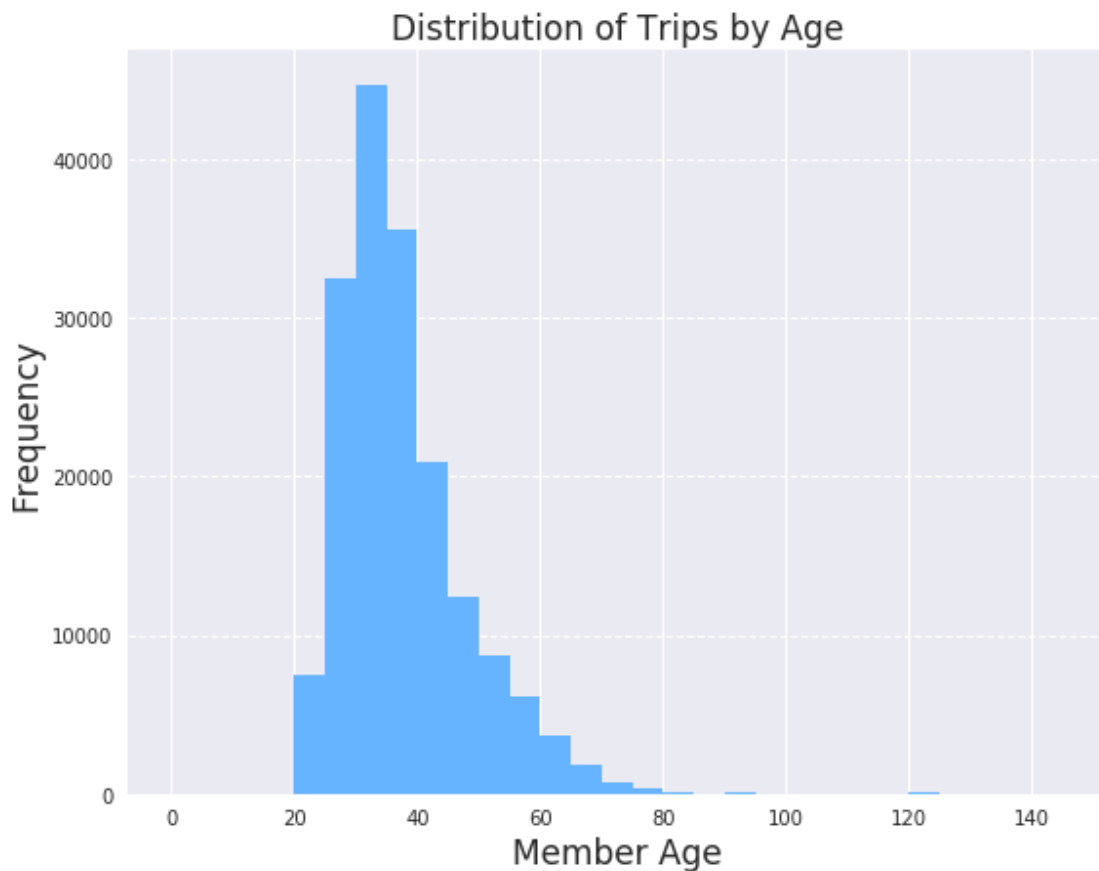


Dips in the trip frequency throughout the month seems to coincide with the weekends.

4.5 Visualization 5: What is the distribution of trips by member age?

```
In [93]: bins = np.arange(0, bike_df['age'].max()+5, 5)
plt.hist(data=bike_df, x='age', bins=bins, color='#66b3ff')

# set title of the histogram
plt.title('Distribution of Trips by Age', fontsize=17)
# set x_label of plot
plt.xlabel('Member Age', fontsize=17)
# set y_label of plot
plt.ylabel('Frequency', fontsize=17)
# add gridlines
plt.grid(axis='y', linestyle='--')
```

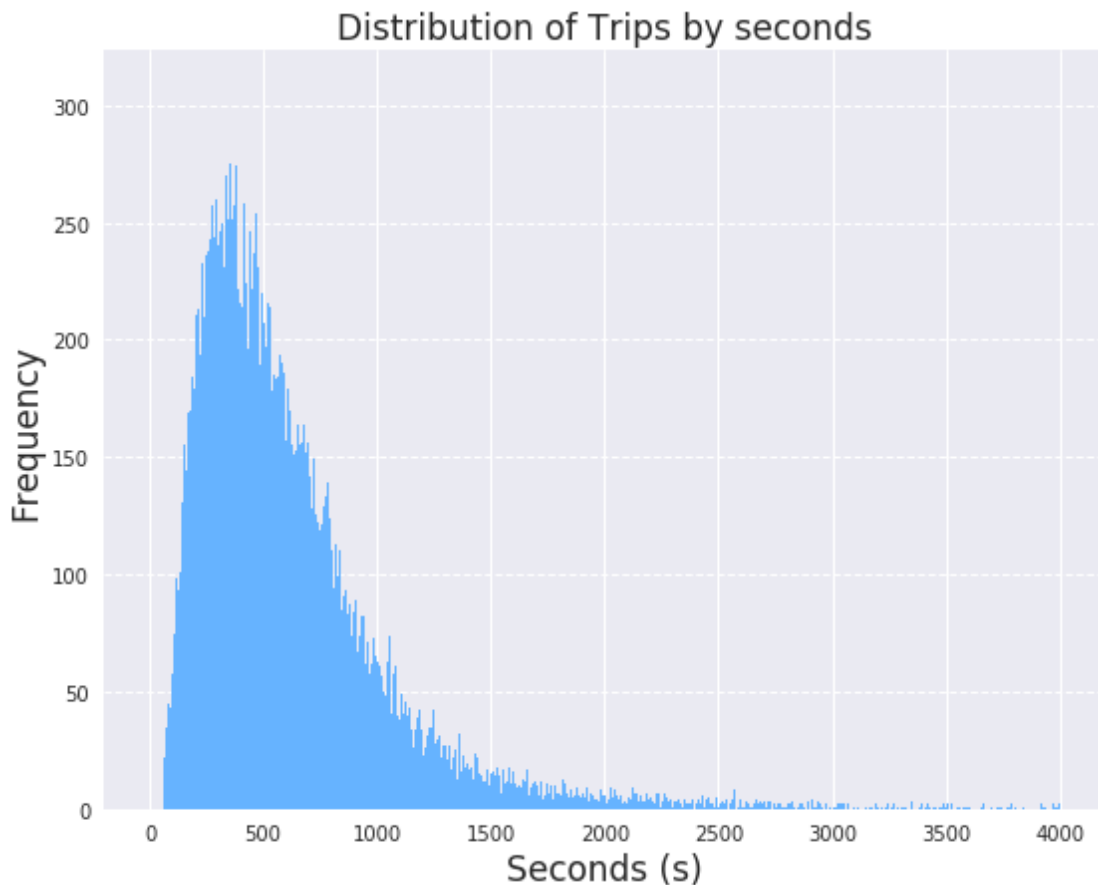


Fewer individuals in their 20s, but interestingly there seems to be datapoints for users at age > 100. Might be an error.

4.6 Visualization 6: What is the distribution of trip by duration (seconds)?

```
In [96]: bins = np.arange(0, 4000, 1)
        ticks = np.arange(0, 100, 5)
        plt.hist(data=bike_df, x='duration_sec', bins=bins, color='#66b3ff')

        # set title of the histogram
        plt.title('Distribution of Trips by seconds', fontsize=17)
        # set x_label of plot
        plt.xlabel('Seconds (s)', fontsize=17)
        # set y_label of plot
        plt.ylabel('Frequency', fontsize=17)
        # add gridlines
        plt.grid(axis='y', linestyle='--')
```



From the histogram, it seems that the distribution of trips duration is right skewed. This tells us most trips are concentrated around the 500-700 seconds duration (8~11 minutes long). This might indicate that individuals typically use the bikes for specific use cases rather than leisure.

4.6.1 Discuss the distribution(s) of your variable(s) of interest. Were there any unusual points? Did you need to perform any transformations?

Trips typically peaked around 8am and 5pm, which seems to be the rush hour time period for working/schooling individuals. There were also significantly more subscribers (90+%) than customers in our dataset.

Most rides happened during the weekdays (Mon-Fri) and were short in duration (< 10 minutes) per trip. Also, a large proportion of users age were around late 30s to early 40s.

4.6.2 Of the features you investigated, were there any unusual distributions? Did you perform any operations on the data to tidy, adjust, or change the form of the data? If so, why did you do this?


The trip duration seems to have some extreme outliers that skewed the chart, this is something that I will attempt to clean in the bivariate exploration stage when the issue becomes more apparent.

5 Bivariate Exploration

5.1 Visualization 7: What is the age distribution of bike users membership type?

```
In [95]: sns.boxplot(data=bike_df, x='user_type', y='age', color='#66b3ff')
```

```
# set title of the histogram
plt.title('Age distribution of bike users by type', fontsize=17)
# set x_label of plot
plt.xlabel('User Type', fontsize=17)
# set y_label of plot
plt.ylabel('User Age', fontsize=17)
# add gridlines
plt.grid(axis='y', linestyle='--')
```




The median user age and age distribution between Customer and Subscribers are quite similar.

5.2 Visualization 8: Average trip duration by day of the week

In [100]: `sns.barplot(data=bike_df, x='start_day', y='duration_sec', color='#66b3ff')`

```
# set title of the barplot
plt.title('Trip duration by day of week', fontsize=17)
# set x_label of plot
plt.xlabel('Day of week', fontsize=17)
# set y_label of plot
plt.ylabel('Duration in Seconds', fontsize=17)
# add gridlines
plt.grid(axis='y', linestyle='--')
```

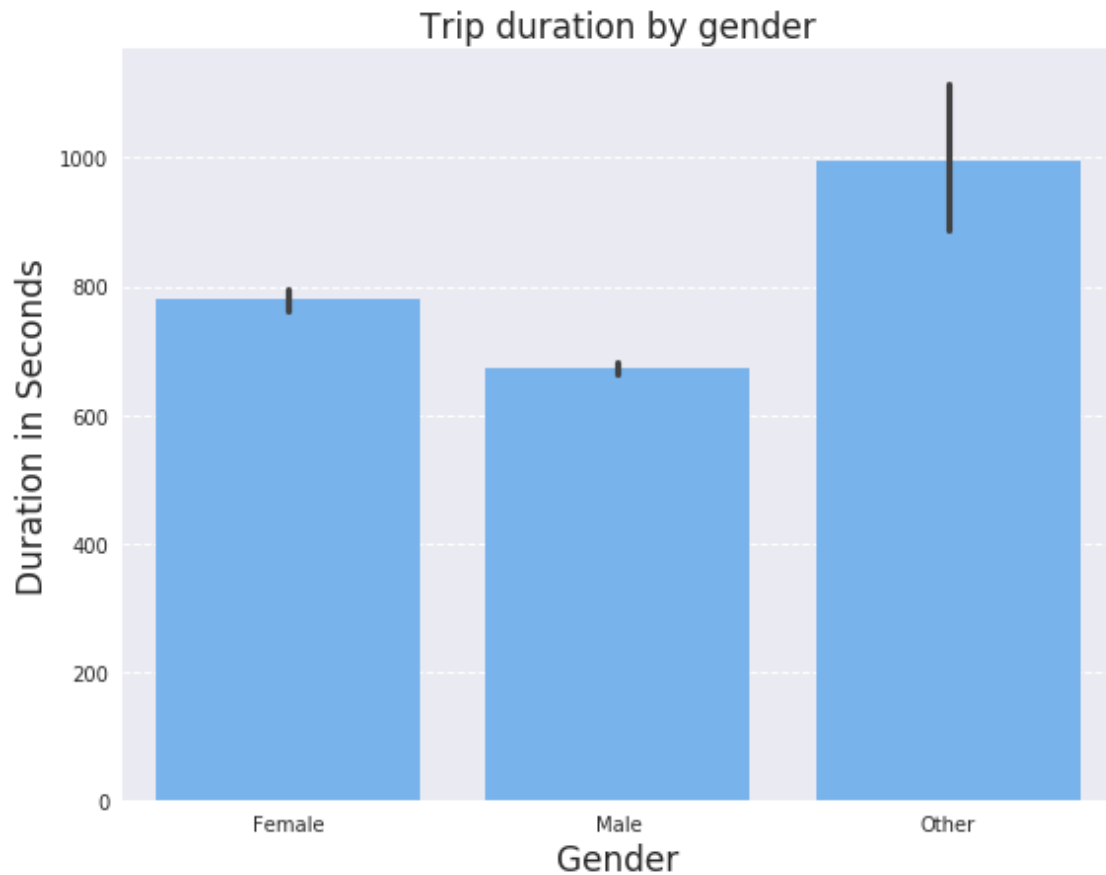


Even though more trips happened during weekdays, the duration of each trip seems to be longer during the weekends. This is very interesting insights.

5.3 Visualization 9: Average trip duration by Gender

```
In [104]: sns.barplot(data=bike_df, x='member_gender', y='duration_sec', color='#66b3ff')
```


```
# set title of the barplot
plt.title('Trip duration by gender', fontsize=17)
# set x_label of plot
plt.xlabel('Gender', fontsize=17)
# set y_label of plot
plt.ylabel('Duration in Seconds', fontsize=17)
# add gridlines
plt.grid(axis='y', linestyle='--')
```

5.4 Visualization 10: Frequency of trips between user types by hour

```
In [110]: sns.countplot(data=bike_df, x='start_hour', hue='user_type')


# set title of the barplot
plt.title('Frequency of trips by User type by hour', fontsize=17)
# set x_label of plot
plt.xlabel('Hour of Day', fontsize=17)
# set y_label of plot
plt.ylabel('Frequency', fontsize=17)
# add gridlines
plt.grid(axis='y', linestyle='--')
```



5.5 Visualization 11: Frequency of trips between user types by day

In [111]: `sns.countplot(data=bike_df, x='start_day', hue='user_type')`


```
# set title of the barplot
plt.title('Frequency of trips by User type by day', fontsize=17)
# set x_label of plot
plt.xlabel('Day of Week', fontsize=17)
# set y_label of plot
plt.ylabel('Frequency', fontsize=17)
# add gridlines
plt.grid(axis='y', linestyle='--')
```



5.6 Visualization 12: Distribution of users' age by Day of the week

```
In [113]: sns.boxplot(data=bike_df, x='start_day', y='age', color='#66b3ff')

# set title of the barplot
plt.title('Distribution of user age by day of the week', fontsize=17)
# set x_label of plot
plt.xlabel('Day of week', fontsize=17)
# set y_label of plot
plt.ylabel('Age', fontsize=17)
# add gridlines
plt.grid(axis='y', linestyle='--')
```




5.7 Visualization 13: Distribution of trip durations by user type

```
In [130]: # creating a minute column
bike_df['duration_min'] = bike_df['duration_sec']/60

# trying to create a violin plot of distribution by minute
sns.violinplot(data=bike_df, x='user_type', y='duration_min', color='#66b3ff', inner='box')

Out[130]: <matplotlib.axes._subplots.AxesSubplot at 0x7f90fb894b00>
```




The duration (in minutes) seems to be extremely skewed by outliers, we will attempt to remove it in the next step.

```
In [132]: # we partition the distribution up till the 99th quantile, allowing us to filter out a
          q = bike_df['duration_min'].quantile(0.99)

          # filtering out for extreme outliers
          duration_outliers_removed = bike_df[bike_df['duration_min'] < q]

          # creating our violin plot
          sns.violinplot(data=duration_outliers_removed, x='user_type', y='duration_min')

          # set title of the barplot
          plt.title('Distribution of user type by minutes', fontsize=17)
          # set x_label of plot
          plt.xlabel('User Type', fontsize=17)
          # set y_label of plot
          plt.ylabel('Duration', fontsize=17)
          # add gridlines
          plt.grid(axis='y', linestyle='--')
```



After plotting our distribution, we can identify the following:

1. Customers had a higher median ride duration as compared to subscribers (as denoted by the white dot in our violinplot).
2. As the "width" of the violin plot represents the kernel density estimation of our dataset, we can conclude that Subscribers have a higher probability of taking shorter trips - as denoted by the thicker sections around the 5min duration range while Customers have a higher probability of taking trips closer to that of 10 minutes.

5.7.1 Talk about some of the relationships you observed in this part of the investigation. How did the feature(s) of interest vary with other features in the dataset?

Subscription use is far higher than customer usage. Customers and subscribers have quite different riding habits and patterns. Most trips on work days (Mon-Fri) and especially during rush hours (when going to work in the morning and getting off work in the afternoon) were made by subscribers because they use the bike sharing system for commuting, whereas customers typically ride for fun in the afternoon or early evening on weekends.

5.7.2 Did you observe any interesting relationships between the other features (not the main feature(s) of interest)?

It's interesting to note that subscribers tend to ride much more quickly and for shorter distances than typical customers. Additionally, subscribers used their bikes the most on weekdays rather than weekends - which signifies that they might be using it for work/school commute rather than for leisure.


5.8 Multivariate Exploration

5.9 Visualization 14: Average trip duration in days of the week between user types

```
In [143]: # creating our pointplot
sns.pointplot(data=bike_df, x='start_day', y='duration_min', hue='user_type', capsize=

# set title of the barplot
plt.title('Comparison of trip duration by days of a week between user types', fontsize=
# set x_label of plot
plt.xlabel('Day of week', fontsize=17)
# set y_label of plot
plt.ylabel('Trip Duration (min)', fontsize=17)
```

```
Out[143]: Text(0,0.5,'Trip Duration (min)')
```



As observed from the point plot, in general: 1. Customers took longer trips as compared to Subscribers 2. There is an increase (though more significant in Customers) on increase in trip duration during the weekends 3. The error bars give us an insight into the measure of central tendency between both user types, whereby customers typically have a much wider spread in trip duration - which is unlike the Subscribers who had a much more concentrated data spread, indicating that users are more likely to have a trip duration close to the median.

This might be due to the terms of bike duration per trip (e.g, Subscribers have free trips for the first xx minutes, thereafter incurring additional charges after the daily free trip duration.)

5.10 Visualization 15: Hourly usage between user types across days of the week

```
In [154]: # defining a function to transform our dataset
def heatmap_transform(df: pd.DataFrame, target:str) -> pd.DataFrame:
    """
    Returns a pivot table of frequency count per day, by the hour

    Args:
        df - A dataframe of n_rows and n_columns
        target - user type we are interested in filtering for
    Returns:
        A pivot table
    """

    transformed_df = (
        df
        # filter out for given target argument
        .query(f'user_type == "{target}"')
        .groupby(['start_day', 'start_hour'])
        .size()
        .reset_index(name='count')
        .pivot(index='start_day', columns='start_hour', values='count')
        .fillna(0)
    )

    return transformed_df

In [156]: cust_count = heatmap_transform(bike_df, 'Customer')
subs_count = heatmap_transform(bike_df, 'Subscriber')

# sample of transformed pivot table
subs_count
```

	00	01	02	03	04	05	06	07	08	09	...	14	15	\
start_day											...			
Monday	86	43	39	20	37	124	556	1545	2966	2062	...	784	1047	
Tuesday	77	51	31	17	30	169	656	2061	3929	3043	...	968	1072	

Wednesday	75	43	29	14	25	152	530	1707	3332	2416	...	859	982
Thursday	109	41	26	19	40	165	661	2121	4167	2569	...	1008	1241
Friday	131	77	55	28	40	138	578	1822	3583	2726	...	786	1089
Saturday	187	145	66	22	19	25	88	188	442	699	...	1076	981
Sunday	137	81	72	33	19	26	50	137	289	519	...	1104	1067

start_hour	16	17	18	19	20	21	22	23
start_day								
Monday	1782	2997	2193	1331	861	561	346	172
Tuesday	2018	3656	2919	1543	1045	622	325	164
Wednesday	1960	3500	2836	1682	1069	815	473	214
Thursday	2357	4065	3300	1800	1282	819	489	277
Friday	1862	2766	2071	1109	621	466	430	256
Saturday	917	930	711	592	395	363	252	201
Sunday	1125	980	825	691	492	415	252	133


[7 rows x 24 columns]

```
In [170]: # defining subplot 1
plt.subplot(2,1,1)
sns.heatmap(cust_count, cmap='rocket_r')
# set x_label of plot 1
plt.xlabel('Day of week', fontsize=10)
# set y_label of plot 1
plt.ylabel('Hour of day', fontsize=10)

# set title of the barplot
plt.title('Hourly trips by day between user types', fontsize='17')

# defining subplot 2
plt.subplot(2,1,2)
sns.heatmap(subs_count, cmap='rocket_r')
# set x_label of plot 2
plt.xlabel('Day of week', fontsize=10)
# set y_label of plot 2
plt.ylabel('Hour of day', fontsize=10)
```

```
Out[170]: Text(60,0.5,'Hour of day')
```



5.10.1 Talk about some of the relationships you observed in this part of the investigation. Were there features that strengthened each other in terms of looking at your feature(s) of interest?

The relationship between the various variables plotted is seen collectively and information is displayed simultaneously, strengthening some of the patterns identified in the preceding bivariate investigation as well as in the univariate research. Subscribers' efficient/short periods of consumption are consistent with their high concentration during Monday through Friday rush hours, showing that the use is mostly for commuting to work. Customers clearly utilize the bike sharing system considerably differently than subscribers, frequently on weekends and in the afternoons, likely for leisure or city tours, as seen by the more flexible and lax pattern of their usage.

5.10.2 Were there any interesting or surprising interactions between features?

When all the interactions between the features are considered together, they all complement one another and make sense, so there aren't any major surprises. Because there are disproportionately more female riders and records than male ones, there may not be a significant difference in usage habits between men and women throughout the

investigation. If there were more female data, it would be interesting to observe the differences in usage between men and women if we were able to obtain a more balanced dataset between both genders

5.11 Conclusions

We performed three levels of visualization and analysis - Univariate, Bivariate and Multivariate on the bike sharing dataset. During the exercise, we also removed extreme outliers to help us better visualize the dataset to generate insights.

Remove all Tips mentioned above, before you convert this notebook to PDF/HTML

At the end of your report, make sure that you export the notebook as an html file from the File > Download as... > HTML or PDF menu. Make sure you keep track of where the exported file goes, so you can put it in the same folder as this notebook for project submission. Also, make sure you remove all of the quote-formatted guide notes like this one before you finish your report!

In []: